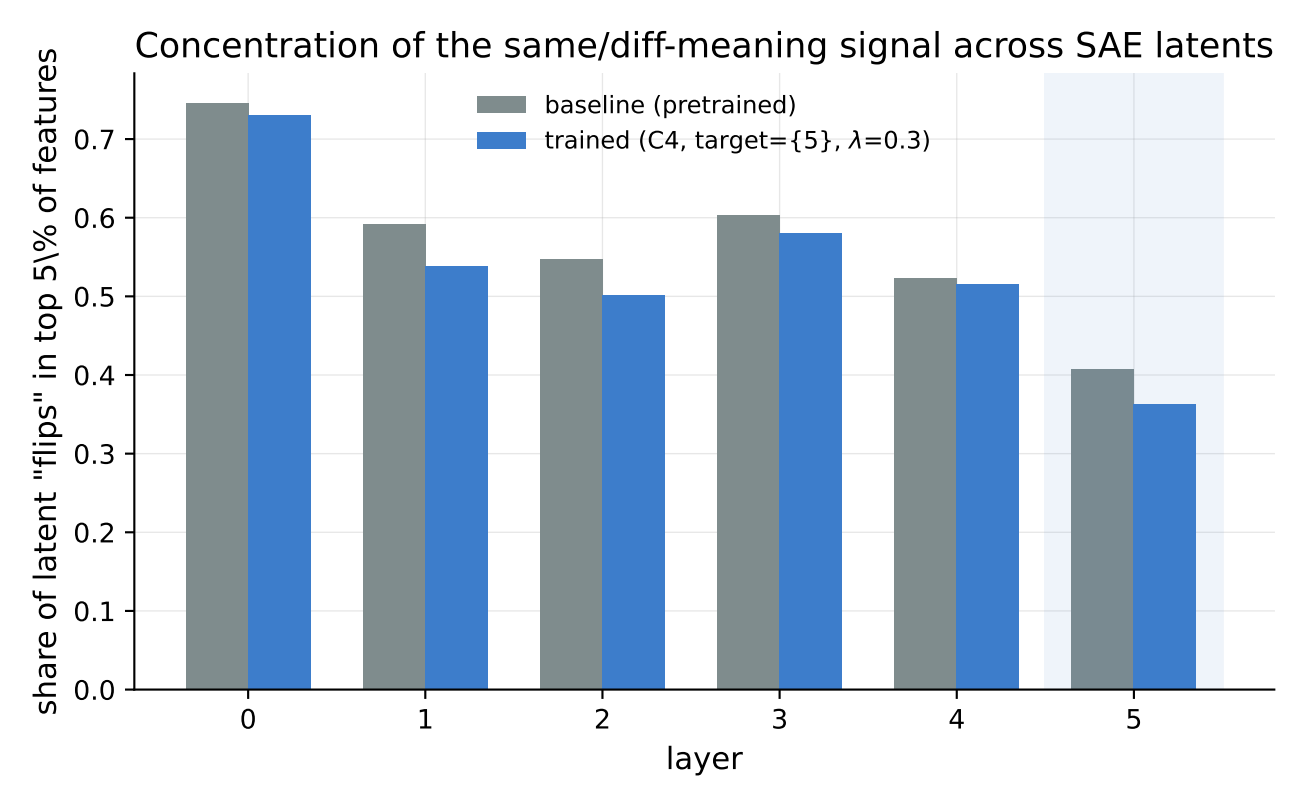
Concentration of the same/diff-meaning signal across SAE latents
baseline (pretrained)
trained (C4, target={5}, λ=0.3)
share of latent "flips" in top 5\% of features
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0
1
2
3
4
5
layer